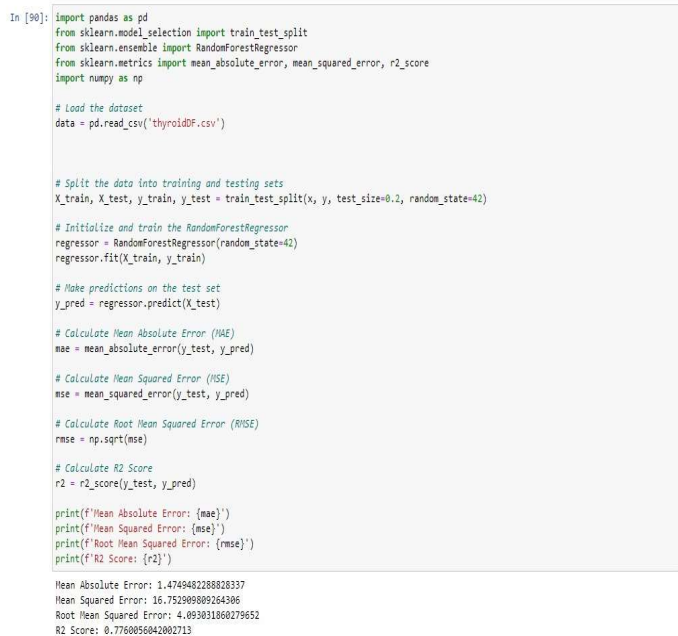


Project Development Phase Model Performance Test

Date	22 November 2023
Team ID	591830
Project Name	Endocrine Elegance: Classifying Thyroid Disorders with Precision
Maximum Marks	10 Marks

Model Performance Testing:

Project team shall fill the following information in model performance testing template.

S.No	Parameter	Values	Screenshot
1.	Metrics	Regression Model: MAE - , MSE - , RMSE - , R2 score - Classification Model: Confusion Matrix - , Accuray Score- & Classification Report -	Regression Model:  Classification Model:

In [92]:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report

# Load the dataset
data = pd.read_csv('thyroid0F.csv')

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)

# Initialize the RandomForestClassifier
classifier = RandomForestClassifier(random_state=42)

# Train the model
classifier.fit(X_train, y_train)

# Make predictions on the test set
y_pred = classifier.predict(X_test)

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print('Confusion Matrix:')
print(conf_matrix)

# Accuracy Score
acc_score = accuracy_score(y_test, y_pred)
print(f'Accuracy Score: {acc_score}')

# Classification Report
class_report = classification_report(y_test, y_pred)
print('Classification Report:')
print(class_report)
```

File Edit View Insert Cell Kernel Widgets Help

Not Trusted

Run

Markdown

Confusion Matrix:

[[1289 3 0 1 0 1 0 1 0 0 10 0 10 0 5
0 0 1 0 1 5 1]
[2 15 0 0 0 0 0 0 0 1 0 0 0 0
0 0 0 0 0 3 0]
[0 1 7 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 2 0]
[4 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0]
[0 0 0 0 28 0 1 0 1 0 0 0 0 0
0 0 0 0 0 0 0]
[0 0 0 0 1 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0]
[0 0 0 0 0 0 69 0 0 0 0 0 0 0
0 0 0 0 0 0 0]
[1 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 6 0 0 0 0 0
0 0 0 0 0 0 0]
[21 3 0 0 0 0 0 0 0 0 57 0 0 0 0
0 0 1 0 0 0 0]
[6 0 0 0 0 0 0 0 0 0 6 0 0 0
0 0 0 0 0 0 0]
[8 0 0 0 0 0 0 0 0 0 0 98 0 0
0 0 0 0 0 0 0]
[1 0 0 0 0 0 0 0 0 0 0 0 1 0
0 0 0 0 0 0 0]
[4 1 0 0 0 0 0 0 0 0 1 0 0 0 15
0 0 0 0 0 1 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0
25 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 6 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 3
0 0 17 0 0 0 0]
[0 0 0 0 2 0 0 0 0 0 0 0 0 0
0 0 0 2 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 1 0 0
0 0 0 2 0 0]
[15 1 1 0 0 0 0 0 0 0 0 1 0 0
0 0 0 0 27 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 22]]

Accuracy Score: 0.9275284359673825

Classification Report:

	precision	recall	f1-score	support
0.0	0.95	0.97	0.96	1328
1.0	0.62	0.71	0.67	21
2.0	0.78	0.78	0.74	10
3.0	0.00	0.00	0.00	4
9.0	0.90	0.95	0.93	40
10.0	0.00	0.00	0.00	1
11.0	0.97	1.00	0.99	69
12.0	0.00	0.00	0.00	1
13.0	0.86	1.00	0.92	6
16.0	0.83	0.70	0.75	82
17.0	1.00	0.50	0.67	12
18.0	0.89	0.92	0.91	106
19.0	1.00	0.50	0.67	2
20.0	0.65	0.54	0.59	28
22.0	1.00	1.00	1.00	25
24.0	1.00	1.00	1.00	6
25.0	0.68	0.85	0.76	20
26.0	1.00	0.50	0.67	4
29.0	0.67	0.67	0.67	3
30.0	0.71	0.60	0.65	45
31.0	0.96	1.00	0.98	22
accuracy			0.93	1835
macro avg	0.74	0.67	0.69	1835
weighted avg	0.92	0.93	0.92	1835

2.	Tune the Model	Hyperparameter Tuning - Validation Method -	<div><div>Hyperparameter Tuning</div><div>Compare the model</div><div><pre>import pandas as pd from sklearn.model_selection import train_test_split, GridSearchCV from sklearn.ensemble import RandomForestClassifier from sklearn.svm import SVC from xgboost import XGBClassifier from sklearn.neural_network import MLPClassifier import matplotlib.pyplot as plt</pre></div><div><pre># Hyperparameter grids param_grids = { 'Random Forest': {'n_estimators': [50, 100, 200], 'max_depth': [None, 10, 20]}, 'SVC': {'C': [1, 10, 100], 'kernel': ['linear', 'rbf']}, 'XGBoost': {'n_estimators': [50, 100, 200], 'max_depth': [3, 6, 9]}, 'ANN': {'hidden_layer_sizes': [(50,), (100,), (50, 50)], 'alpha': [0.0001, 0.001, 0.01]} }</pre></div><div><pre>models = {'Random Forest': rf_model, 'SVC': svc_model, 'XGBoost': xgb_model, 'ANN': ann_model}</pre></div><div><pre>y_train = y_train.astype(int)</pre></div><div><pre># Hyperparameter tuning and training for model_name, model in models.items(): param_grid = param_grids[model_name] grid_search = GridSearchCV(model, param_grid, cv=5, scoring='accuracy', n_jobs=-1)</pre></div><div><pre>grid_search.fit(X_train, y_train)</pre></div><div><pre>C:\Users\Leena\anaconda3\lib\site-packages\sklearn\model_selection_split.py:700: UserWarning: The least populated class as only 1 members, which is less than n_splits=5. warnings.warn(</pre></div><div><div><div>GridSearchCV</div><div><div>estimator: MLPClassifier</div><div>MLPClassifier</div></div></div></div><div><pre>best_model = grid_search.best_estimator_ # Evaluate on training set train_acc = best_model.score(X_train, y_train) accuracy_scores_train.append(train_acc) # Evaluate on test set test_acc = best_model.score(X_test, y_test) accuracy_scores_test.append(test_acc) # Print results for each model print(f'{model_name} - Best Parameters: {grid_search.best_params_}') print(f'{model_name} - Training Accuracy: {train_acc:.4f}') print(f'{model_name} - Test Accuracy: {test_acc:.4f}') print()</pre></div><div><pre>ANN - Best Parameters: {'alpha': 0.001, 'hidden_layer_sizes': (50, 50)} ANN - Training Accuracy: 0.9849 ANN - Test Accuracy: 0.9738</pre></div></div>
----	----------------	---	--

			<div><pre>plt.figure(figsize=(3, 6)) plt.bar(model_name, accuracy_scores_train, labels='Training Set', alpha=0.7) plt.bar(model_name, accuracy_scores_test, labels='Testing Set', alpha=0.7) plt.xlabel('Models') plt.ylabel('Accuracy') plt.title('Model Accuracy Comparison') plt.legend() plt.show()</pre></div> <div><p>Model Accuracy Comparison</p><table><tr><th>Model</th><th>Training Set Accuracy</th><th>Testing Set Accuracy</th></tr><tr><td>ANN</td><td>1.0</td><td>0.95</td></tr></table></div>	Model	Training Set Accuracy	Testing Set Accuracy	ANN	1.0	0.95
Model	Training Set Accuracy	Testing Set Accuracy							
ANN	1.0	0.95							
			<div>Validation Method</div> <div><pre>In [91]: import pandas as pd from sklearn.model_selection import cross_val_score, train_test_split from sklearn.ensemble import RandomForestRegressor from sklearn.metrics import make_scorer, mean_squared_error import numpy as np # Load the dataset data = pd.read_csv('thyroidDF.csv') # Split the data into training and testing sets X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42) # Initialize the RandomForestRegressor regressor = RandomForestRegressor(random_state=42) # Define a custom scoring function (you can choose the metric you're interested in) scorer = make_scorer(mean_squared_error, greater_is_better=False) # Perform cross-validation (e.g., using 5-fold cross-validation) cv_scores = cross_val_score(regressor, X_train, y_train, cv=5, scoring=scorer) # Calculate mean and standard deviation of the cross-validation scores mean_cv_score = np.mean(cv_scores) std_cv_score = np.std(cv_scores) print(f'Mean Cross-Validation Score: {mean_cv_score}') print(f'Standard Deviation of Cross-Validation Scores: {std_cv_score}') # Train the model on the full training set regressor.fit(X_train, y_train) # Evaluate the model on the test set y_pred = regressor.predict(X_test) test_score = mean_squared_error(y_test, y_pred) print(f'Mean Squared Error on Test Set: {test_score}')</pre><div>Mean Cross-Validation Score: -13.465419040376939 Standard Deviation of Cross-Validation Scores: 0.8274871413731842 Mean Squared Error on Test Set: 16.75290809264306</div></div>						